

DISTRIBUTED COMPUTING AND MAPREDUCE

CKME 134 – BIG DATA ANALYTICS TOOLS

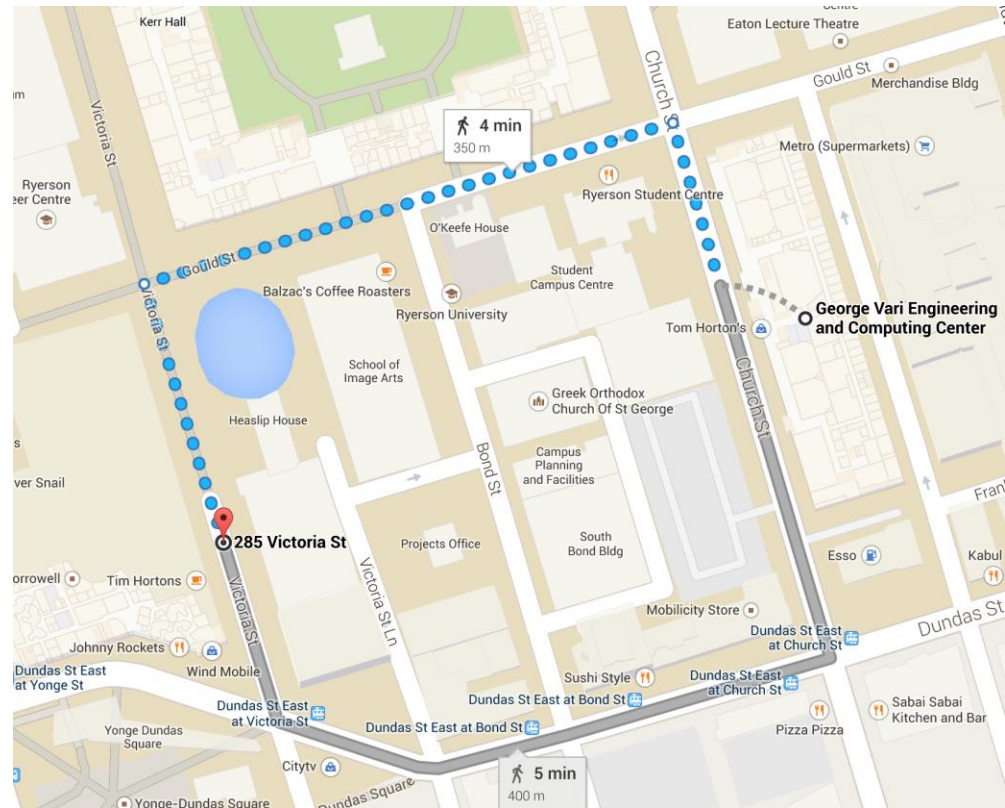
RYERSON UNIVERSITY

SPRING 2015

Instructor: Shaohua Zhang

General Course Information

- Instructor
 - ▣ Shaohua Zhang
 - ▣ Ryerson shaohua.zhang@ryerson.ca
 - ▣ Personal shaohua.zhang@live.com
- GA
 - ▣ Behjat Soltanifar
 - ▣ behjat.soltanifar@ryerson.ca
- Lectures
 - ▣ 6:30~8:30
 - ▣ ENGLG06
- Lab
 - ▣ 8:30~9:30
 - ▣ 285 Victoria St (403/404)
 - Take the elevator to 4FL



Course Outline (*subject to change*)

1. Intro to Big Data
2. **Distributed Computing and MapReduce**
3. Hadoop Ecosystem
4. Intro to Hive
5. Pig
6. Advanced Pig
7. Hadoop Performance Optimization
8. Big Data Use Cases: Location Intelligence and Marketing Analytics
9. Big Data Use Cases: Recommendation Engine and Computational Advertising
10. Hadoop In Action: Building Data Pipelines
11. Beyond Hadoop: Spark
12. Beyond Hadoop: Real-Time Analytics

Assignment Schedule

<i>Date</i>	<i>Out</i>	<i>In</i>
<i>Assignment 1</i>	<i>Jan 12</i>	<i>No Due Date</i>
<i>Assignment 2</i>	<i>Hive</i>	
<i>Assignment 3</i>	<i>Pig</i>	
<i>Assignment 4</i>	<i>Data Pipeline</i>	



Lecture 1 Recap

Course Blackboard

Lecture 1 Review

Watson and future of analytics

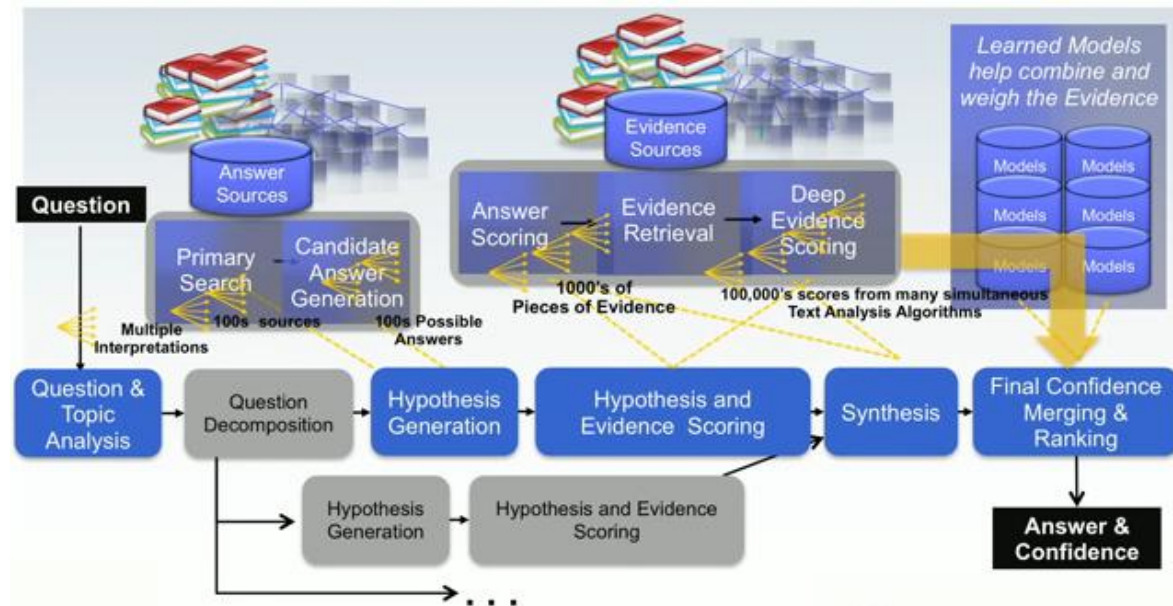
Questions from students

Lecture 1 Recap - Review

1. Big Data Introduction
2. Big Data Use Cases
3. Data Analytics Tooling
4. Big Data Job Market
5. Big Data Challenges
 - ▣ [file:///localhost/Users/DSinmotion/Dropbox/Startup/RyersonToolsCourse/Winter 2015/Session1-Introduction-to-big-data-analytics-tools.pptx](file:///localhost/Users/DSinmotion/Dropbox/Startup/RyersonToolsCourse/Winter%202015/Session1-Introduction-to-big-data-analytics-tools.pptx) - 43. 4. Big Data Challenges

Session 1 Recap – Watson and Future of Analytics

- DeepQA Project – The Jeopardy Game
- Future of analytics
 - ▣ Computers will replace human to some extent
 - Watson
 - Deep Learning
 - ▣ Machine learning automation
 - Enterprise Miner, KXEN etc.
 - Vincent Granville → consulting (analytics automation)
- We're still in the early days
 - ▣ Human judgment is still critical
 - ▣ We need many more Watson like projects to make it real



Building Watson: An Overview of the DeepQA Project <http://www.aaai.org/ojs/index.php/aimagazine/article/download/2303/2165>

Building Watson: A Brief Overview of the DeepQA Project (youtube) <https://www.youtube.com/watch?v=3G2H3DZ8rNc>

Your cognitive future: How next-gen computing changes the way we live and work: [http://www-](http://www-935.ibm.com/services/us/gbs/thoughtleadership/cognitivefuture/)

[935.ibm.com/services/us/gbs/thoughtleadership/cognitivefuture/](http://www-935.ibm.com/services/us/gbs/thoughtleadership/cognitivefuture/)

Session 1 Recap - Questions From You

- Students with advanced tools knowledge
 - ▣ Start building your data science portfolio
 - Build your github repo
 - Join open source projects
 - Kaggle – a place to polish your machine learning skills and learn from other great data scientists
 - Build something real

Session 1 Recap - Questions From You

- Students business/marketing background who are NOT interested in pursuing data scientist career
 - ▣ Being good with data is becoming increasingly important
 - Most MBA programs are opening data analytics courses
 - Orgs are becoming more and more data driven
 - ▣ Communities are building user-friendly big data platforms for you
 - SQL, scripting languages are your friends
 - ▣ Understanding basic stats/data mining concepts are very helpful when communicating with the data team

Big Data Use Case: News & Media

- Newspaper (Globe and Mail, NYT, Bloomberg)
 - ▣ Content curation → editor vs. computer
 - Which piece of news should make the headline?
 - Content classification (topic modeling)
 - ▣ Lead generation
 - Where do you get new digital subscribers from?
 - Online advertising → Profile your existing user base and then bid on those user attributes
 - Traditional media → Where/when do you place your ads?
 - ▣ Revenue
 - Subscription model
 - Advertising – publisher revenue optimization
 - ▣ User tracking/understanding
 - How do you track your readers?
 - Browser session based
 - Email subscription
 - User demographic prediction
 - ▣ Content personalization
 - News recommendation
- Big data
 - Hadoop, text mining, topic modeling, recommendation etc.

Big Data Use Case: News & Media

□ Bloomberg

▣ News Recommender

- Big data engineer (*15k ~ 20k salary*)
- http://www.slideshare.net/Hadoop_Summit/shah-june27-425pmroom210av2
- <http://www.youtube.com/watch?v=nNAbBXc1EYo>

□ New York Times

▣ Content digitization

- The New York Times used Amazon's EC2 compute cloud to crunch through four terabytes of scanned archives from the paper, converting them to PDFs for the Web

Lecture 2 - Outline


- Distributed Computing
- MapReduce
- Algorithms
 - ▣ Word Count
 - ▣ K-Means
- Beyond MapReduce



Distributed Computing

Single Node Architecture

- Traditionally, computation has been CPU bound
 - ▣ Complex computation on small data
- For decades, the primary push is to increase the computing power of a single machine




Scale Up vs. Scale Out

- Single Node Architecture
 - ▣ Scaling up advantage
 - Programming is easier than distributed computing
 - Faster processing on smaller data
 - ▣ Scale up disadvantage
 - Hardware cost
 - Scalability
- Advantage of scale-out systems
 - ▣ Scalability
 - ▣ Cost

Traditional Distributed Systems: Problems


- Problems with traditional distributed systems:
 - ▣ Complex programming model
 - ▣ Network bandwidth is the bottleneck
 - ▣ It is difficult to deal with partial failures of the system
 - ▣ Typically at compute time, data is copied to the compute nodes
 - This doesn't scale to today's big data problems!



Data Becomes the Bottleneck

- Traditional distributed systems don't scale to today's Internet-scale data
- Getting data to the computer processor becomes the bottleneck
 - ▣ Disk I/O is slow
 - ▣ Network bandwidth is bottleneck
- Solution → moving computation to the data!

<ul style="list-style-type: none">• <i>Internet</i>	<ul style="list-style-type: none">○ 2.5 exabytes (2.5×10^{18}) per day – 2012○ 2.3 zettabytes (2.3×10^{21}) per day - 2014
<ul style="list-style-type: none">• <i>Facebook</i>	<ul style="list-style-type: none">○ 500+ terabytes per day○ 100+ petabytes in a single Hadoop cluster

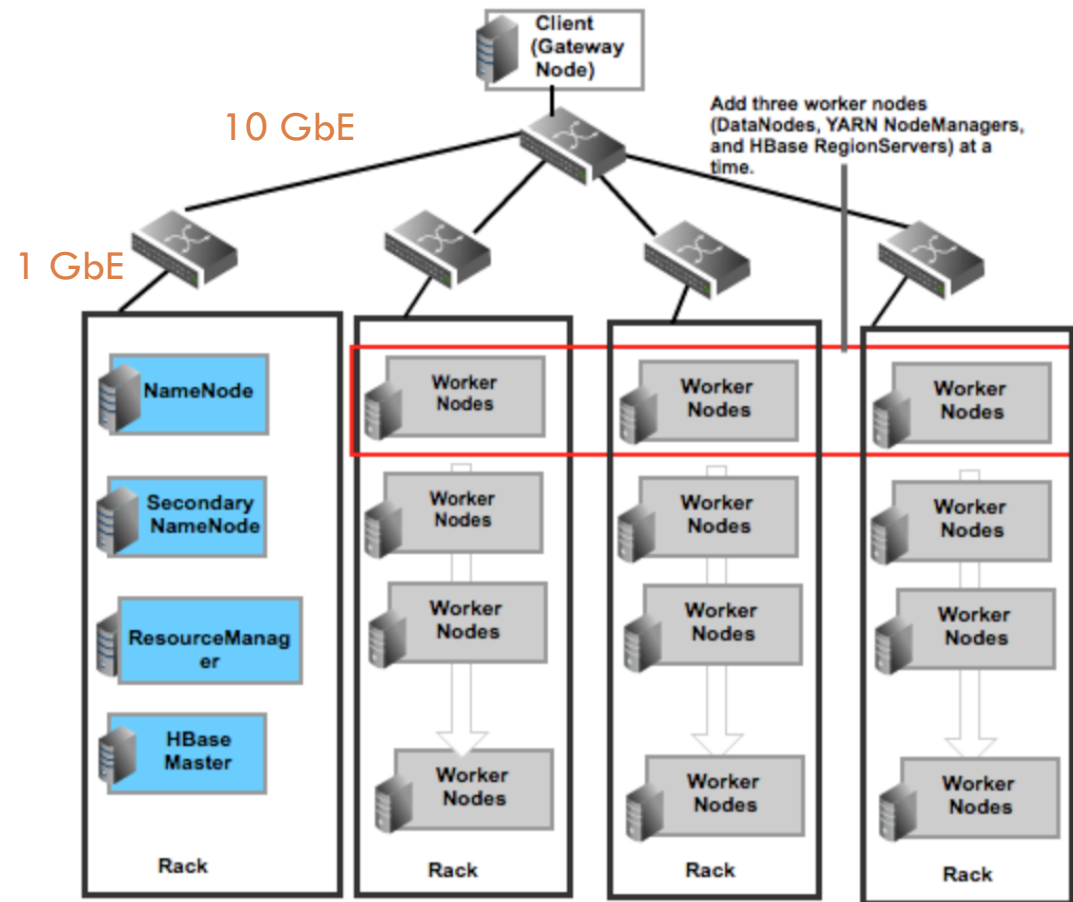


MapReduce to the rescue!

Modern Distributed Computing Cluster

Cluster architecture

- A medium-to -large Hadoop cluster consists of a two-level or three-level architecture built with rack-mounted servers. Each rack of servers is interconnected using a 1 Gigabyte Ethernet switch. Each rack-level switch is connected to a cluster-level switch (which is typically a larger port-density 10GbE switch).



NOTE: DataNodes, NodeManagers, and RegionServers are typically co-deployed.

Big Data Made Possible


□ Hardware

▣ Big cluster of commodity machines at lower cost

- Faster processor
- Cheaper memory
- Bigger hard drive space
- Faster network bandwidth

□ Software

▣ Algorithms to allow parallel computing (map-reduce)





MapReduce

Setting the Expectation

- If your goal is to become a data analyst, you probably don't have to learn MapReduce programming
 - ▣ Alternatively, you need to be good with Pig/Hive
 - ▣ Still, it is important to understand the M/R basics
- But if you want to become a Hadoop data architect, data engineer or research engineer or maybe data scientist who work with very large data...
 - ▣ You'll need to understand M/R programming patterns well
 - ▣ You'll still use Pig/Hive 80% of your time, but being able to use M/R allows you to do more complex data processing
 - Writing customized UDF (user-defined functions) in Pig/Hive also requires a good understanding of M/R


MapReduce

- *MapReduce is a computing model that decomposes large data manipulation jobs into individual tasks that can be executed in parallel across a cluster of servers*
- Each node processes data stored on that node
- Consists of two phases
 - ▣ Map
 - ▣ Reduce

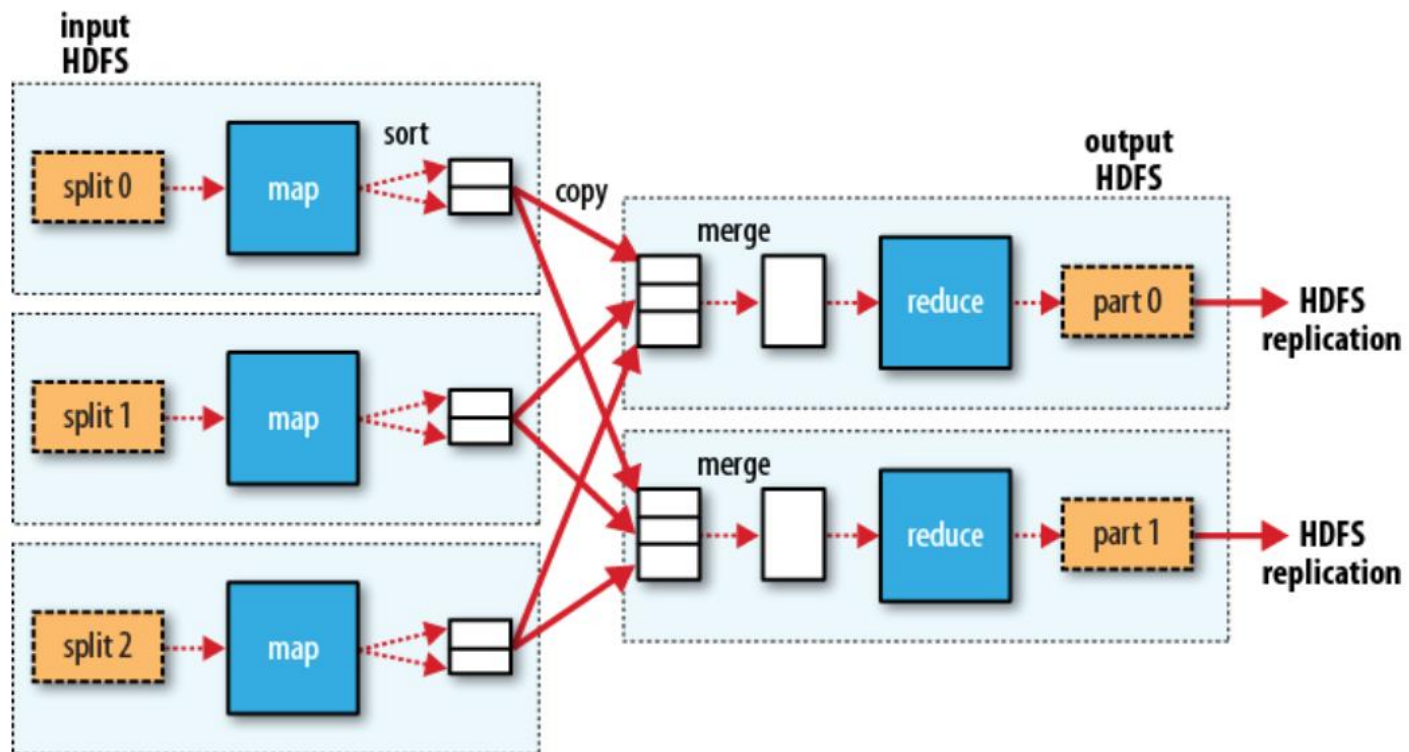
MapReduce

- Automatic parallelization and distribution
 - ▣ It makes M/R programming much easier
- Developer simply need to focus on writing the *map* and *reduce* functions
- M/R is written in Java
- It also supports Python Streaming
 - ▣ writing *map* and *reduce* function in python

MapReduce – Toy Example



MapReduce – Map & Reduce



Word Count Example

- Word Count is a classic programming example, used in many tutorials
- Counting word occurrences is widely used in many natural language processing related tasks
 - ▣ TF-IDF (Term Frequency-Inverse Document Frequency) is key input into many complex algorithms such as PageRank and document classification

Note: we will walk through TF-IDF in more detail in later lectures

Word Count Explained

- The WordCount program reads/scans through the document line by line
- It tokenizes/splits the line by delimiters (space, tab etc.)
- Each occurrence of a word/term will increment the corresponding word count by 1

I	wish	to	wish	the
wish	you	wish	to	
wish	but	if	you	
wish	the	wish	the	
witch	wishes	,	I	
won't	wish	the	wish	
you	wish	to	wish	

Document

I	---->	1 1
wish	---->	1 1 1 1 1 1 1 1 1 1 1 1
to	---->	1 1 1
the	---->	1 1 1 1
you	---->	1 1 1
but	---->	1
if	---->	1
witch	---->	1
wishes	---->	1
won't	---->	1

I	---->	2
wish	---->	1 1
to	---->	3
the	---->	4
you	---->	3
but	---->	1
if	---->	1
witch	---->	1
wishes	---->	1
won't	---->	1

MapReduce handles these automatically for you!!



```
package org.myorg;
```

```
import java.io.IOException;
import java.util.*;
```

```
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
```

```
public class WordCount {
```

```
    public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();
```

```
        public void map(LongWritable key, Text value, Context context)
            throws IOException, InterruptedException {
            String line = value.toString();
            StringTokenizer tokenizer = new StringTokenizer(line);
            while (tokenizer.hasMoreTokens()) {
                word.set(tokenizer.nextToken());
                context.write(word, one);
            }
        }
    }
}
```

```
    public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {
```

```
        public void reduce(Text key, Iterable<IntWritable> values, Context context)
            throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
        }
    }
}
```

```
    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
```

```
        Job job = new Job(conf, "wordcount");
```

```
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
```

```
        job.setMapperClass(Map.class);
        job.setReducerClass(Reduce.class);
```

```
        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);
```

```
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

```
        job.waitForCompletion(true);
    }
}
```

MapReduce

Hive

```
CREATE TABLE docs (line STRING);
```

```
LOAD DATA INPATH 'docs' OVERWRITE INTO TABLE docs;
```


```
CREATE TABLE word_counts AS
SELECT word, count(1) AS count FROM
  (SELECT explode(split(line, '\s')) AS word FROM docs) w
GROUP BY word
ORDER BY word;
```

K-Means Clustering

- Clustering Use Cases
 - ▣ Customer value segmentation
 - High value customer
 - ▣ Behavior segmentation
 - Marketing
 - Sales
 - Product
 - ▣ Location clustering
 - ▣ Multi-stage algorithms
 - Stage 1 → clustering
 - Stage 2 → classification


K-Means – Single Node

- Kmeans → Iterative algorithm until convergence
 1. Select K points at random as cluster centroids (centers)
 - Now we formed K clusters
 2. For each data point, assign it to the closest center
 3. For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
 4. If the new centers are different from the old centers (previous iteration) → Go to Step 2
 - Otherwise, stop




K-Means – Single Node

- Kmeans → Iterative algorithm until convergence
 1. Select K points at random as cluster centroids (centers)
 2. For each data point, assign it to the closest center
 - Now we formed K clusters
 3. For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
 4. If the new centers are different from the old centers (previous iteration) → Go to Step 2
 - Otherwise, stop




K-Means – Single Node

- Kmeans → Iterative algorithm until convergence
 1. Select K points at random as cluster centroids (centers)
 - Now we formed K clusters
 2. For each data point, assign it to the closest center
 - Now we formed K clusters
 3. For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
 4. If the new centers are different from the old centers (previous iteration) → Go to Step 2
 - Otherwise, stop




K-Means – Single Node

- Kmeans → Iterative algorithm until convergence
 1. Select K points at random as cluster centroids (centers)
 - Now we formed K clusters
 2. For each data point, assign it to the closest center
 - Now we formed K clusters
 3. For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
 4. If the new centers are different from the old centers (previous iteration) → Go to Step 2
 - Otherwise, stop



K-Means – Single Node

- Kmeans → Iterative algorithm until convergence
 1. Select K points at random as cluster centroids (centers)
 2. For each data point, assign it to the closest center
 - Now we formed K clusters
 3. For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
 4. If the new centers are significantly different from the old centers (previous iteration)
 - Set the new centers then Go to Step 2
 - Otherwise, stop




K-Means – Single Node

□ Kmeans → Iterative algorithm until convergence

1. Select K points at random as cluster centroids (centers)
2. For each data point, assign it to the closest cluster
 - Now we have formed K clusters
3. For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
4. If the new centers are significantly different from the old centers (previous iteration)
 - Set the new centers then Go to Step 2
 - Otherwise, stop


Iteration 2 Starts

Iteration 2




K-Means – Single Node

- Kmeans → Iterative algorithm until convergence
 1. Select K points at random as cluster centroids (centers)
 2. For each data point, assign it to the closest center
 - Now we formed K clusters
 3. For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
 4. If the new centers are significantly different from the old centers (previous iteration)
 - Set the new centers then Go to Step 2
 - Otherwise, stop



K-Means – Single Node


- Kmeans → Iterative algorithm until convergence
 1. Select K points at random as cluster centroids (centers)
 2. For each data point, assign it to the closest center
 - Now we formed K clusters
 3. For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
 4. If the new centers are significantly different from the old centers (previous iteration)
 - Set the new centers then Go to Step 2
 - Otherwise, stop



K-Means – Single Node

- Kmeans → Iterative algorithm until convergence
 1. Select K points at random as cluster centroids (centers)
 - 2. For each data point, assign it to the closest center
 - Now we formed K clusters
 3. For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
 4. If the new centers are significantly different from the old centers (previous iteration)
 - Set the new centers then Go to Step 2
 - Otherwise, stop

Iteration 2




K-Means – Single Node

- Kmeans → Iterative algorithm until convergence


1. Select K points at random as cluster centroids (centers)
2. For each data point, assign it to the closest cluster
 - Now we have formed K clusters
3. For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
4. If the new centers are significantly different from the old centers (previous iteration)
 - Set the new centers then Go to Step 2
 - Otherwise, stop

Iteration 3




K-Means – Single Node

- Kmeans → Iterative algorithm until convergence
 1. Select K points at random as cluster centroids (centers)
 2. For each data point, assign it to the closest center
 - Now we formed K clusters
 3. For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
 4. If the new centers are significantly different from the old centers (previous iteration)
 - Set the new centers then Go to Step 2
 - Otherwise, stop



K-Means – Single Node

- Kmeans → Iterative algorithm until convergence
 1. Select K points at random as cluster centroids (centers)
 2. For each data point, assign it to the closest center
 - Now we formed K clusters
 3. For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
 4. If the new centers are significantly different from the old centers (previous iteration)
 - Set the new centers then Go to Step 2
 - Otherwise, stop




K-Means – Single Node

- Kmeans → Iterative algorithm until convergence
 1. Select K points at random as cluster centroids (centers)
 2. For each data point, assign it to the closest center
 - Now we formed K clusters
 3. For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
 4. If the new centers are significantly different from the old centers (previous iteration)
 - Set the new centers then Go to Step 2
 - Otherwise, stop

Meets convergence criteria


Iteration 3



K-Means – Single Node

- Kmeans → Iterative algorithm until convergence
 1. Select K points at random as cluster centroids (centers)
 2. For each data point, assign it to the closest center
 - Now we formed K clusters
 3. For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
 4. If the new centers are significantly different from the old centers (previous iteration)
 - Set the new centers then Go to Step 2
 - Otherwise, stop

Meets convergence criteria



Distributed K-Means - MapReduce

□ Map-side

- Each map reads the K-centroids + one block from dataset
- Assign each point to the closest centroid
- Output <centroid, point>

□ Reduce-side

- Gets all points for a given centroid
- Re-compute a new centroid for this cluster
- Output: <new centroid>

□ Iteration Control

- Compare the old and new set of K-centroids
 - If similar → Stop
 - Else
 - If max iterations has reached → Stop
 - Else → Start another Map-Reduce Iteration

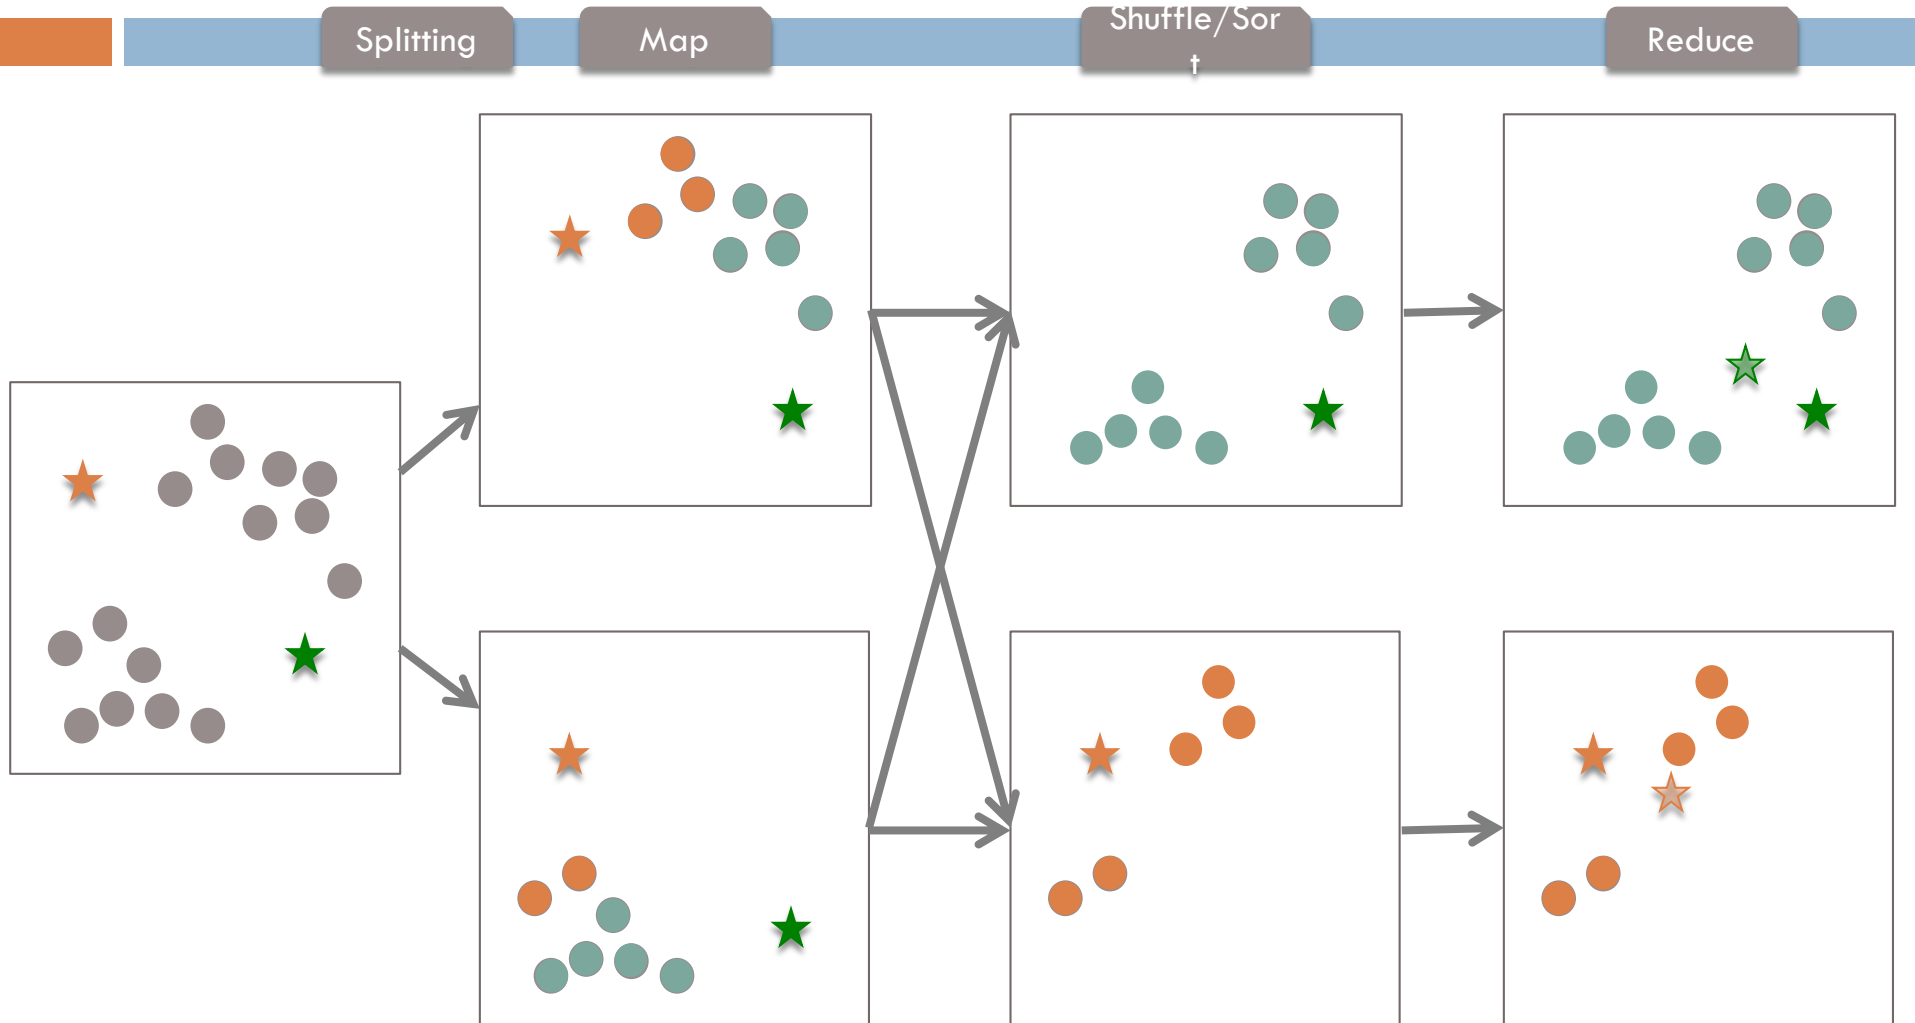
□ Use of combiners

- Similar to the reducer
- Computes for each centroid the local sums (and counts) of the assigned points
- Sends to the reducer <centroid, <partial sums>>

□ Use of single reducer

- Amount of data to reducers is very small
- Single reducer can tell whether any of the centers has changed or not
- Creates a single output file

Distributed K-Means - MapReduce



K-Means M/R One Iteration




Recommended Readings - MapReduce

- (Google) Google MapReduce ([link](#))
- (Google) The Google File System ([link](#))
- (Microsoft) Scaling up vs scale out for hadoop: time to rethink? ([link](#))
- (Google) Vision paper: towards an understanding of the limits of mapreduce computation ([link](#))
- (Twitter) MapReduce is good enough? If all you have is a hammer, throw away everything that is not a nail! ([link](#))

Recommended Readings – MapReduce/Hadoop

□ Books




- ▣ Hadoop: The Definitive Guide – Chap 2
- ▣ Mining of Massive Datasets – Chap 2.1,2.2
- ▣ Data-Intensive Text Processing with MapReduce – Chap 2



Recommended Readings – Data Analytics

□ The Data Analytics Handbook

▣ <https://www.teamleada.com/handbook#download>



Recommended Reading – Linux Commands

- Data Science at the Command Line

- ▣ <http://datascienceatthecommandline.com>

