

ADVANCED PIG

CKME 134 – BIG DATA ANALYTICS TOOLS

RYERSON UNIVERSITY

SPRING 2015

Instructor: Shaohua Zhang


Course Outline

2

1. Intro to Big Data
2. Distributed Computing and MapReduce
3. Hadoop Ecosystem
4. Programming Hive
5. Advanced Hive
6. Mid-Term Review
7. Programming Pig
8. **Advanced Pig**
9. Hadoop Performance Optimization
10. Hadoop In Action: Building Data Pipelines
9. Location Analytics and Recommender Systems
11. Beyond Hadoop: Spark
12. Beyond Hadoop: Graph Analytics

Data Scientist?

3




Josh Wills

@josh_wills

+ Follow

Data Scientist (n.): Person who is better at statistics than any software engineer and better at software engineering than any statistician.



RETWEETS

988

FAVORITES

486



9:55 AM - 3 May 2012

(video) Airbnb Tech Talk: Josh Wills - The Life of a Data Scientist

<https://www.youtube.com/watch?v=h9vQIPfe2uU>

Recap

4

Type	Description	Example
<i>int, long, float, duple, chararray, bytearray, boolean, datetime</i>	<ul style="list-style-type: none">• <i>primitive data types in pig</i>	
<i>tuple</i>	<ul style="list-style-type: none">• <i>Fixed length, ordered set of fields, like a row with multiple columns in SQL</i>• <i>Allows random access</i>• <i>Must fit in memory</i>	<i>(toronto, 3)</i> <i>('bob', 53, 'toronto', 'male')</i>
<i>bag</i>	<ul style="list-style-type: none">• <i>An unordered collection of tuples</i>• <i>Can have tuples with differing numbers of fields</i>• <i>Can spill to disk (doesn't have to fit in memory)</i>	<i>{(toronto, 3),(chicago, 5)}</i> <i>{('bob',53,'toronto','male'), ('sally', 23), 'george', 'montreal')}</i>
<i>map</i>	<ul style="list-style-type: none">• <i>A set of key/value pairs</i>• <i>Key/values in a relation must be unique</i>• <i>Key must be chararray, data element can be any type</i>	<i>[city#toronto]</i> <i>[name#bob,age#53,city#toronto,gender#male]</i>

Tuple

5

□ Tuple

`twitter:tuple(id:chararray, lat:double, lon:double, tweet:chararray)`

`twitter.id` → `USER_12345678`

`twitter.$3` → `#Hadoop, big data is the new oil`

`twitter.($1,$2)` → `(40.48956, -156.22234)`

`twitter.(lat,lon)` → `(40.48956, -156.22234)`

Bag

6

<i>id</i>	<i>field1</i>	<i>field2</i>	<i>field3</i>
<i>user1</i>	<i>a</i>	<i>b</i>	<i>c</i>
<i>user2</i>	<i>a</i>	<i>b</i>	
<i>user3</i>	<i>a</i>		

```
grunt> data = load '/user/lab/pig/data_test_bag' using PigStorage('\t') as  
(id:chararray, f1:chararray, f2:chararray, f3:chararray);  
grunt> grpd = group data ALL;
```

```
grunt> describe grpd;
```

```
grpd: {group: chararray,data: {(id: chararray,f1: chararray,f2: chararray,f3: chararray)}}
```

```
grunt> dump grpd;
```

```
(all,{(user3, a, , ), (user2, a, b, ), (user1, a, b, c)})
```

```
cnt_id = foreach grpd generate COUNT(data.id) as cnt; ← 3
```

```
cnt_f1 = foreach grpd generate COUNT(data.$1) as cnt; ← 3
```

```
cnt_f2 = foreach grpd generate COUNT(data.f2) as cnt; ← 2
```

```
cnt_f3 = foreach grpd generate COUNT(data.$3) as cnt; ← 1
```

Complex Type

7

- Type Construction Operators/Functions
 - ▣ Tuple construction $(a, b) \rightarrow$ Same as TOTUPLE()
 - ▣ Bag construction $\{a, b\} \rightarrow$ Same as TOBAG()
 - ▣ Map construction $[a, b] \rightarrow$ Same as TOMAP()

Bag Construction

8

```
A = load 'students' as (name:chararray, age:int, gpa:float);  
B = foreach A generate {(name, age)}, {(name, age)};  
store B into 'results';
```

Input (students):

joe smith	20	3.5
amy chen	22	3.2
leo allen	18	2.1

Output (results):

```
{(joe smith,20)} {(joe smith), (20)}  
{(amy chen,22)} {(amy chen), (22)}  
{(leo allen,18)} {(leo allen), (18)}
```


Map Construction

9

```
A = load 'students' as (name:chararray, age:int, gpa:float);  
B = foreach A generate [name, gpa];  
store B into 'results';
```

Input (students):

```
joe smith 20 3.5  
amy chen 22 3.2  
leo allen 18 2.1
```

Output (results):

```
[joe smith#3.5]  
[amy chen#3.2]  
[leo allen#2.1]
```

Comparison Operator

10

```
data = load '/user/lab/pig/full_text.txt' AS (id:chararray,  
ts:chararray, location:chararray, lat:float, lon:float,  
tweet:chararray);
```

```
filtr = FILTER data BY tweet MATCHES '.*@USER_\\S{8}.*';
```

```
limt = limit filtr 20;
```

```
dump limt;
```

Scalar Projections

11

- Pig allows you to cast the elements of a single-tuple relation into a scalar value
- The primary use case for casting relations to scalars is the ability to use the values of global aggregates in follow up computations

Scalar Projection Demo

Normalize average number of tweets per user

12

```
a = load '/user/lab/pig/full_text.txt' AS (id:chararray, ts:chararray,  
location:chararray, lat:float, lon:float, tweet:chararray);
```

```
b = group a by id;
```

```
c = foreach b generate group as id, COUNT(a) as user_cnt;
```

```
d = group c ALL;
```

```
e = foreach d generate AVG(c.user_cnt) as global_avg;
```

```
f = foreach c generate id, user_cnt/(float)e.global_avg as index;
```

```
store f into '/user/lab/pig/tweet_count_index';
```

id	cnt
user1	45
user2	22
user3	67

	global_avg	
ALL	44.67	
id	cnt	index
user1	45	1.00
user2	22	0.49
user3	67	1.50

COGROUP

13

- COGROUP is a generalization of GROUP that works with more relations
- You can COGROUP up to but no more than 127 relations at a time
- The result is a record with a key and one bag for each input. Each bag contains all records from that input that have the given value for the key
- Another way to think of COGROUP is as the first half of a join. The keys are collected together, but the cross product is not done
- COGROUP is useful when you want to do join-like things but not a full join
- $\text{COGROUP} + \text{FLATTEN} = \text{JOIN}$

COGROUP Demo

uid	age	gender	region
u1	14	M	US
u2	32	F	UK
u3	22	M	US

uid	region
u1	US
u1	UK
u1	CA
u2	US

14

```
grunt> user = load '/user/lab/pig/user.txt' using PigStorage(',') as (uid:chararray,  
age:int, gender:chararray, region:chararray);  
grunt> session = load '/user/lab/pig/session.txt' using PigStorage(',') as  
(uid:chararray, region:chararray);
```

```
grunt> cogrp = COGROUP user BY uid, session BY uid;  
grunt> describe cogrp;  
cogrp: {group: chararray,user: {(uid: chararray,age: int,gender: chararray,region:  
chararray)},session: {(uid: chararray,region: chararray)}}}
```

```
grunt> dump cogrp;  
(u1,{(u1,14, M, US)}, {(u1, CA),(u1, UK),(u1, US)})  
(u2,{(u2,32, F, UK)}, {(u2, US)})  
(u3,{(u3,22, M, US)}, {})
```

```
grunt> cogrp_nest = foreach cogrp {  
  crossed = cross user, session;  
  generate crossed;  
}  
grunt> dump cogrp_nest;  
(({u1,14, M, US, u1, CA}, {u1,14, M, US, u1, UK}, {u1,14, M, US, u1, US}))  
(({u2,32, F, UK,u2, US}))
```

Set Intersection

with COGROUP

name	hits	name	errors
John	3	John	2
Harry	4	John	3
George	2	George	0
		Sue	1

15

```
grunt> s1 = load 'hits-data' as (name:chararray, hits:int);  
grunt> s2 = load 'errors-data' as (name:chararray, errors:int);
```

```
grunt> grps = COGROUP s1 BY name, s2 BY name;
```

```
(John, {(John, 3)}, {(John, 2), (John, 3)})  
(Harry, {(Harry, 4)})  
(George, {(George, 2), (George, 2)})  
(Sue, {}, {(Sue, 1)})
```

-- Note:

-- Something is in the intersection of s1 and s2 if there are no {}'s in the cgroup.

```
grunt> grps2 = FILTER grps by NOT(IsEmpty(s1)) AND NOT(IsEmpty(s2));  
grunt> intersection = FOREACH grp2 GENERATE group as grp, s1, s2 ;  
grunt> dump intersection;
```

```
(John, {(John, 3)}, {(John, 3), (John, 2)})  
(George, {(George, 2)}, {(George, 0)})
```

Set Difference

with COGROUP

name	hits	name	errors
John	3	John	2
Harry	4	John	3
George	2	George	0
		Sue	1

16

```
grunt> s1 = load 'hits-data' as (name:chararray, hits:int);  
grunt> s2 = load 'errors-data' as (name:chararray, errors:int);
```

```
grunt> grps = COGROUP s1 BY name, s2 BY name;
```

```
(John, {(John, 3)}, {(John, 2), (John, 3)})  
(Harry, {(Harry, 4)})  
(George, {(George, 2), (George, 2)})  
(Sue, {}, {(Sue, 1)})
```

-- Note:

-- Something is in the difference between s1 and s2 if there are non-empty second sets.

```
grunt> grps2 = FILTER grps by !isEmpty(s2);  
grunt> set_diff = FOREACH grps2 GENERATE group as grp, s1, s2 ;  
grunt> dump set_diff;
```

```
(Harry, {(Harry, 4)}, {})
```


Advanced Joins

17

- replicated join
- skewed join
- semi-join (co-group)
- non-equi join (cross)

Replicated Join

18

- ❑ Joins small data to large data
- ❑ Often used when a lookup file is in a smaller file
 - ▣ Small enough to fit in memory of your computer node
- ❑ Using **distributed cache**
- ❑ Replicated join applies map-only join
 - ▣ No reduce phase – Yeaah!!
- ❑ Supports only inner and left outer join
- ❑ Can be used with more than two tables.
- ❑ In this case, all but the first (left-most) table are read into memory
 - ▣ The lookup file should always be on the right side

Replicated JOIN Demo

Finding Weekend Tweets

19

```
a = load '/user/lab/pig/full_text.txt' using PigStorage('\t') AS (id:chararray,  
ts:chararray, location:chararray, lat:float, lon:float, tweet:chararray);
```

```
a1 = foreach a generate id, ts, SUBSTRING(ts,0,10) as date;
```

```
b = load '/user/lab/pig/dayofweek.txt' using PigStorage('\t') as  
(date:chararray, dow:chararray);
```

```
b1 = filter b by dow=='Saturday' or dow=='Sunday';
```

```
c = join a1 by date, b1 by date using 'replicated';
```

```
d = foreach c generate a1::id .. a1::date, b1::dow as dow;
```

```
e = limit d 5;
```

```
dump e;
```

Skewed Join `c = join a by key, b by key using 'skewed';`

20

- Skewed join is used when the data has significant skew in the number of records per key. This results in one or two reducers that will take much longer than the rest → Straggler problem
- Adds 5% overhead to the performance due to the sampling approach at the beginning
- Skew join works by first sampling one input for the join. In that input it identifies any keys that have so many records that skew join estimates it will not be able to fit them all into memory.
- Those keys identified as too large are treated differently. Based on how many records were seen for a given key, those records are split across the appropriate number of reducers. The number of reducers is chosen based on Pig's estimate of how wide the data must be split such that each reducer can fit its split into memory. For the input to the join that is not split, those keys that were split are then replicated to each reducer that contains that key.

Non-Equi Join (CROSS)

find user with more than 2 friends

p_name	value	name	friend
nfriends	2	Amy	George
ndays	100	George	Fred
nvists	13	Fred	Anne
		George	Joe
		George	Harry

21

```
grunt> params = load '/user/lab/pig/params.txt' using PigStorage(',') as
(p_name:chararray, value:int);
grunt> friend = load '/user/lab/pig/friend.txt' using PigStorage(',') as (name:chararray,
friend:chararray);
grunt> friend_grp = group friend by name;
grunt> friend_cnt = foreach friend_grp generate group as name, COUNT(friend.friend)
as cnt;

grunt> friend_param = filter params by p_name=='nfriends';
grunt> friend_param_p = foreach friend_param generate value;

grunt> friend_cross = CROSS friend_cnt, friend_param_p;

grunt> friend_cross_1 = filter friend_cross by friend_cnt::cnt >= friend_param_p::value;
grunt> dump friend_cross_1;
```

name	friend
Amy	1
George	3
Fred	1

name	friend	fr_param_p
Amy	1	2
George	3	2
Fred	1	2

George	3	2
--------	---	---

Sampling

22

-- approach 1

```
grunt> data = load 'file';  
grunt> sample1 = SAMPLE data 0.1;
```

-- approach 2

```
grunt> data = load 'file';  
grunt> sample2 = FILTER A by random() <= 0.1;
```

-- approach 3

```
grunt> data = load 'file';  
grunt> grpd = group a all;  
grunt> grpd_cnt = foreach b generate COUNT(data) as num_rows;  
grunt> sample3 = SAMPLE a 10000/grpd_cnt .num_rows;
```

UDF – User Defined Function

23

- Java
 - ▣ Java has the most extensive support. You can customize all parts of the processing including:
- Jython/Python/Javascript/Ruby/Groovy
 - ▣ Limited support is provided for Python, Jython etc.
- Where to get UDFs?
 - ▣ Built-In Functions
 - ▣ Piggybank/DataFu/Pigeon/etc.
 - Download from project site (maven central)
 - Download source and compile

Git & Github

24

- Git Installation and GitHub Setup
 - ▣ Refer to “Git & GitHub” section in Module 8 BlackBoard

UDF Usage

25

- Find the JAR online via maven central
 - ▣ OR download source and build your own in Eclipse
- Upload the JAR to Hadoop Sandbox
- In Pig
 - ▣ Register the JAR first
 - ▣ Define the functions to use
 - ▣ Invoke the function

Piggybank

<http://search.maven.org>

26

-- piggybank UDFs

```
REGISTER '/home/lab/piggybank-0.14.0.jar';  
DEFINE isotounix org.apache.pig.piggybank.evaluation.datetime.convert.ISOToUnix();  
  
a = load '/user/lab/pig/full_text.txt' AS (id:chararray, ts:chararray, location:chararray,  
lat:float, lon:float, tweet:chararray);  
b = foreach a generate id, ts, isotounix(ts) as ts_unix;  
c = limit b 3;  
dump c;
```

DataFu – Random Sampling

27

-- Simple Random Sampling (SRS)

-- Take a 1% random sample from the dataset

-- Register UDFs and define functions first

register /home/lab/datafu-1.2.0.jar

DEFINE SRS datafu.pig.sampling.SimpleRandomSample('0.01');

-- Simple Random Sampling (SRS)

data = load '/user/lab/pig/full_text.txt' AS (id:chararray, ts:chararray, location:chararray, lat:float, lon:float, tweet:chararray);

sampled = foreach (group data all) generate flatten(**SRS**(data));

store sampled into '/user/lab/pig/full_text_src';

-- Stratified Sampling (by date)

-- Take a 1% random sample from each date using SRS and group by date

data = load '/user/lab/pig/full_text.txt' AS (id:chararray, ts:chararray, location:chararray, lat:float, lon:float, tweet:chararray);

data1 = foreach data generate id, SUBSTRING(ts, 0, 10) as date, lat, lon, tweet;

grouped = **GROUP** data1 **BY** date;

sampled = foreach grouped generate flatten(**SRS**(data1));

store sampled into '/user/lab/pig/full_text_stratified';

Pigeon – UDF for GeoSpatial

28

- Download and build pigeon jars
 - ▣ git clone <https://github.com/aseldawy/pigeon.git>
 - ▣ import project as maven in eclipse
 - ▣ build jar
 - ▣ under “target” sub-folder, find jar file
 - ▣ Upload to sandbox via ftp
- Demo

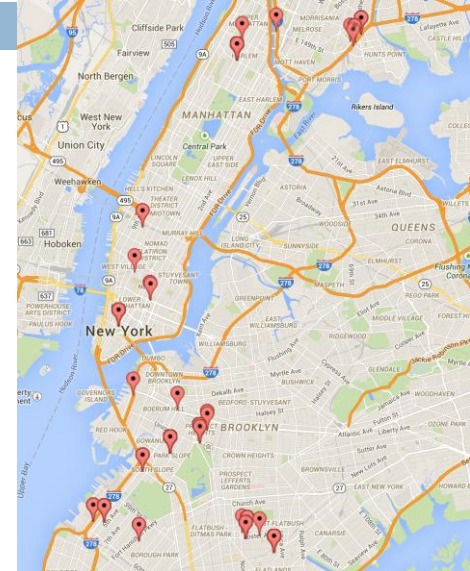
Pigeon UDF Demo

Find tweets tweeted from NYC

29

```
-- register UDFs
REGISTER /home/lab/pigeon-1.0-SNAPSHOT.jar;
REGISTER /home/lab/esri-geometry-api-1.2.jar;

-- define functions
DEFINE ST_MakeBox edu.umn.cs.pigeon.MakeBox;
DEFINE ST_Contains edu.umn.cs.pigeon.Contains;
DEFINE ST_MakePoint edu.umn.cs.pigeon.MakePoint;
```



```
data = load '/user/lab/pig/full_text.txt' AS (id:chararray, ts:chararray,
location:chararray, lat:double, lon:double, tweet:chararray);
data1 = FOREACH data GENERATE id, ts, lat, lon, ST_MakePoint(lat, lon)
AS geom_point, tweet;
data2 = FILTER data1 BY ST_Contains(ST_MakeBox(40.4774, -74.2589,
40.9176, -73.7004), geom_point);
data3 = limit data2 200;
data4 = foreach data3 generate lat, lon;
dump data4;
```

<http://www.darrinward.com/lat-long/?id=490564>